

LEDGER-Q delivery checklist

6 October 2026. This records implemented behavior and its limits. It is not a claim that the product is qualified for commercial regulated use.

Working product

F01 — Controlled Quality Source & Document Registry

Permitted TXT/PDF/image originals, version/hash/predecessor, checked candidate transcription; PDF automation remains external. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F02 — Requirement Extraction & Exact Evidence Traceability

Exact text requirement candidates with source/version/span/offset; named review controls acceptance. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F03 — Regulation-to-SOP / Requirement Mapping

SOP coverage, rationale, linked requirement and source records; missing support blocks approval. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F04 — Compliance Gap, Conflict & Evidence-Sufficiency Control

Saved gaps, conflict and dependency states; unsupported candidates fail grounding checks. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F05 — Quality Risk, Severity & Confidence Assessment

Separate severity, likelihood, controls and confidence fields; human root cause and disposition. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F06 — Deviation Investigation Review

Investigation chronology, hypotheses, human conclusions and immutable contribution trail. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F07 — OOS/OOT Investigation Review

Synthetic OOS assay conflict 94.2% vs 99.1%; missing rationale remains HOLD. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F08 — Bidirectional Change-Control & Impact Assessment

Transitive saved links, quality-to-AI invalidation, AI-to-quality links and historical decisions preserved. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F09 — RCA-CAPA Lifecycle & Effectiveness Verification

Implementation/effectiveness/cause gates, due dates and signed closure; missing proof blocks closure. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F10 — Training Impact & Retraining Traceability

Role/version/completion evidence and fresh review requirements. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F11 — Audit / Inspection Finding & Readiness Management

Finding, owner, linked action and evidence packet; no formal inspection-readiness certification. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F12 — PQR/APQR & Quality Trend Review

Saved counts, denominators and time window; no statistically proven quality trend claim. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F13 — Owner, Deadline, Escalation & Closure Control

Owner/due date, visible overdue work and saved manual escalation; no external delivery configured. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F14 — Human QA Review, Override & Controlled Correction

Current role/revision/hash, password plus real TOTP, decisions and reviewed regression proposals. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F15 — Audit/Version Reconstruction, Event Forensics & Controlled Export

Revision replay, append audit, R2 packets, checksummed backup and isolated copy restore. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F16 — AI Model / Use / Context & Release Registry

Intended/excluded use, model/config/expiry, signed passport states; synthetic authority only. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F17 — AI Provenance, Transparency & Contribution Record

Saved parent/specialist/reviewer outputs, attempts, input hashes, correlation and human contribution boundary. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F18 — AI Evaluation, Grounding & Hallucination Suite

Executed deterministic controls plus signed frozen cases and optional provider comparison. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F19 — Bias, Edge-Case & Consistency Verification

Explicit subgroups and two-model repeated checks; live inference and representative fairness evidence not established. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

F20 — Assurance Monitoring, Regression, Blast Radius & Revalidation

Change/expiry invalidation, visible assurance debt, approved correction proposals and fresh revalidation/review. — tests/integration.mjs + saved UI records; provider paths use mock contract tests unless a workspace configures live access

Nine agentic parts

Goal / planning

Saved bounded parent plan, evidence checks and one held-case next plan; no regulatory release planning.

Tools

Scoped D1/R2 source, extraction, provider and export adapters; fixed Groq endpoint and document consent.

Memory

Current workspace records, prior snapshots, frozen cases and reviewed corrections.

Orchestration

Saved phases/checkpoints, leases, idempotency, pause/resume/cancel and exhausted/manual-recovery state. A run executes through a request; no independently scheduled worker queue is claimed.

Multi-agent collaboration

Seven named bounded roles; candidate checks and different-model independent review when configured. No unrestricted swarm.

Human review

Explicit synthetic reviewer role plus password/TOTP, exact revision/hash and rationale.

Safety

Tenant/role/source-purpose limits, file checks, untrusted source isolation and blocked human-authority actions.

Observability

Actual saved attempts, errors, run duration, available provider token usage, first failures and evaluation denominators. No fabricated cost or speed benefit.

Governance

Original provenance, revision audit, signature records, correction firewall, expiry and controlled export. Retention policy metadata exists; automated retention deletion and regulatory certification are not claimed.

Design and navigation

Landing → inspect saved proof, or create a separate synthetic workspace → Quality workspace / AI assurance / Review and evidence.

Source → check → human review → saved record. Each evidence panel has close controls and a clear original-opening confirmation.

Approved ice-blue/slate palette, six distinct bento cards, compact expandable navigation, scroll reveal and pointer tilt.

Silver/glass hero is generated raster artwork with continuous CSS motion; it is not an articulated WebGL simulation. Reduced-motion preferences and hidden-tab pauses remain supported.

The three supplied videos and four approved design previews remain in the complete package as reference material, not as captured footage of the implementation.

Tests and limits

39 D1-compatible SQLite / R2-memory integration checks pass. Paired model response tests are mocked; no paid inference or professional outcomes are inferred.

TypeScript check and official Worker build must pass before deployment. Final native deployment result is recorded separately.

Browser visual QA was unavailable because the required control-browser capability was absent. Responsive rules and accessible modal/navigation primitives are implemented, but pixel-perfect browser verification is not claimed.

Workspace aggregate is bounded to 1,500 records / 1.8 MB current state. Backup limits: 2 MB referenced original bytes and 4 MB packet. This is not globally scaled enterprise storage.

Backup copy retains originals and the imported historical bundle; it starts a new local review history. Imported decisions cannot grant current approval.

What needs real access or professional evidence

Configure this workspace's own Groq key, two accessible model IDs, optional vision model and processing permission in Controls. No existing project keys were copied.

Automated PDF/table parsing, email verification/recovery, Google OAuth, external notification delivery and unattended queue scheduling require selected real services and further integration.

Representative held-out data, qualified reviewer answers, supervised users, measured manual baseline, production load/security assessment and full disaster-recovery drills remain necessary before paying regulated users.

No real QA qualification, commercial CMC experience, Part 11 certification, autonomous regulatory release or percentage improvement is claimed.